

PhiGraph Core 4.0: A Shadow-First Evidence Ledger and Policy-Gated Runtime for Software-Agent Operations

Walter Calmels von Dem Knesebeck
TUCH Systems, Santiago, Chile
wcalmels@phi47.cl

August 2026 (v2 draft)

Abstract

Software agents and artificial-intelligence systems can produce claims, recommendations, and action proposals faster than organizations can verify and govern them. This paper presents PhiGraph Core 4.0.0, an MIT-licensed, model-agnostic system that treats model output as a proposal rather than as verified truth. PhiGraph represents claims, evidence, verifications, action proposals, policy decisions, and outcomes as linked records; stores them in a scope-tagged, tamper-evident ledger; and places policy evaluation between agent proposals and operational execution. Replay and shadow modes cannot execute external actions, while higher-authority modes require both an explicit executor and an allowing policy decision. We describe the protocol, runtime, ledger, deployment boundaries, and Python implementation. Core 4.1 adds a scoped transactional ledger with declared write locks and fail-closed chain verification for JSON and SQLite backends. Building on this substrate, we additionally describe PhiGraph HAV v0.2, a fail-closed verification layer that extracts claims from agent-generated text, checks them against an explicitly supplied authoritative state, and records the resulting claims, evidence, verifications, and policy decisions in the same ledger. Engineering verification comprises 128 automated tests (117 core tests and 11 HAV-specific tests); the continuous-integration configuration targets Python 3.10–3.13. An initial external experiment on the CIC-IDS2017 Friday DDoS flow file compares a PhiGraph relational score with three unsupervised baselines. Local Outlier Factor is strongest on this tabular benchmark; PhiGraph ranks second by PR-AUC and attains 0.840 precision at a 10% alert budget. The experiment demonstrates functional integration, not state-of-the-art superiority: it uses a single imperfect dataset, a feature-similarity graph, and one fixed split. PhiGraph is therefore presented as an auditable governance substrate and research prototype, not as a proof of correctness, causal identification, or safe autonomous operation.

Keywords: AI governance; software agents; provenance; evidence ledger; graph intelligence; policy enforcement; shadow deployment; claim verification; reproducible software

1 Introduction

Modern software agents can synthesize text, modify code, recommend operational changes, invoke tools, and coordinate multi-step workflows. Their outputs may be useful without being reliable, authorized, or supported by reproducible evidence. This creates a control problem: an organization must determine what was claimed, which evidence supported or refuted it, which policy governed a proposed action, who approved it, and what happened after evaluation or execution.

Existing guidance emphasizes lifecycle risk management and traceability. The NIST AI Risk Management Framework organizes activities around governing, mapping, measuring, and managing

AI risk [10]. W3C PROV provides a general vocabulary for provenance [5], while model cards and dataset datasheets improve documentation of released artifacts [7, 2]. These approaches are important, but they do not by themselves define an operational boundary between a model-generated proposal and an authorized action.

PhiGraph addresses this boundary through an evidence-oriented protocol and a policy-gated runtime. Its central invariant is that model output is a candidate claim or action proposal, not verified truth or implicit authority. The system records relationships among claims, evidence, verification results, policies, approvals, and outcomes so that a decision can be reconstructed after the fact.

This paper makes five contributions:

1. a typed protocol for claims, evidence, verifications, action proposals, policy decisions, and outcomes;
2. a runtime in which policy evaluation and an explicit executor are necessary conditions for external action;
3. a backend-neutral, scope-tagged ledger with hash-chain integrity metadata and optional HMAC evidence signatures;
4. HAV v0.2, a fail-closed extension that verifies agent-generated textual claims against an authoritative external state and records the result through the same protocol and ledger; and
5. a bounded empirical evaluation comprising software verification and an external anomaly-ranking experiment, together with explicit threats to validity.

2 Related Work

2.1 Governance, documentation, and provenance

The NIST AI RMF is a voluntary, use-case-agnostic framework for managing AI risk [10]. PhiGraph is not an implementation or certification of that framework; it provides operational records that may support governance activities. PROV-O formalizes entities, activities, agents, and provenance relations [5]. PhiGraph uses a narrower application protocol centered on verification and policy-gated action. A future interoperability layer could map PhiGraph records to PROV-O.

Model cards [7] and datasheets for datasets [2] make intended uses, evaluation conditions, and limitations explicit. PhiGraph applies a related transparency principle at runtime: each claim and action remains linked to its evidence, decision, and outcome rather than relying only on static release documentation.

2.2 Relational representations and software assurance

Graph-based representations make entities and relationships first-class and support structured reasoning [1]. PhiGraph uses heterogeneous graphs and classical graph analytics; it does not require a graph neural network and does not claim that graph structure alone establishes causality. In software assurance, in-toto protects software-supply-chain integrity through verifiable step metadata [12]. PhiGraph shares an emphasis on attributable records but targets a broader claim–evidence–policy–action lifecycle.

Agentic systems introduce risks including tool misuse, privilege compromise, untraceable actions, and cascading errors. The OWASP Agentic Security Initiative provides an industry threat taxonomy

[8]. PhiGraph’s deny-by-default policy, approval requirements, idempotency controls, and shadow-first deployment are engineering controls relevant to these risks; they are not a complete security solution.

2.3 Automated claim verification

Large language models can generate fluent text that is not supported by, or directly contradicts, available evidence, a phenomenon surveyed broadly across natural-language-generation tasks [3]. FEVER established claim verification against textual evidence as a distinct research problem, using retrieval and entailment models over a large Wikipedia-derived claim set [11]. PhiGraph’s HAV extension (Section 4.6) targets a narrower, operational instance of this problem: claims that an AI agent makes about the status of a software system, checked against a caller-supplied authoritative state rather than an open-domain knowledge base. It does not use a trained entailment model and does not claim comparable coverage or accuracy to FEVER-style systems; its contribution is a fail-closed policy binding between extracted claims and PhiGraph’s existing evidence ledger and governance protocol.

3 System Model and Design Goals

Let an agent produce a proposal

$$P = (C, A),$$

where C is a set of claims and A is a set of proposed actions. A claim does not become verified merely because an agent assigns it high confidence. Instead, a verifier records a result over a claim and an explicit evidence set. Similarly, an action is not authorized by its presence in P : the policy engine evaluates the action under a runtime mode and an approval set.

PhiGraph is designed around six goals:

1. **Evidence boundedness:** when evidence identifiers are supplied, verification results reference evidence already registered in the ledger.
2. **Least authority:** absence of a matching policy blocks an action.
3. **Traceability:** claims, decisions, and outcomes remain reconstructable.
4. **Safe evaluation:** replay and shadow runs never execute an external action.
5. **Scope tagging:** tenant and project identifiers propagate through stored records and can be used to filter queries.
6. **Model neutrality:** deterministic programs, statistical models, and language-model agents can use the same protocol.

4 PhiGraph Architecture




Figure 1: PhiGraph separates agent proposals from verification, policy evaluation, and execution. Arrows indicate recorded processing relationships, not automatic trust propagation.

4.1 Canonical protocol

Protocol version 2.0.0 exposes six principal frozen dataclass record schemas:

- **Claim:** statement, type, subject, issuer, status, confidence, evidence references, and optional

supersession;

- **Evidence**: kind, source, payload, status, content hash, observation time, and metadata;
- **Verification**: claim, verifier, method, result, evidence references, rationale, and time;
- **ActionProposal**: action type, target, proposer, parameters, rationale claims, reversibility, and risk level;
- **PolicyDecision**: effect, applicable policies, reasons, required approvals, and time; and
- **Outcome**: action, status, whether execution occurred, observed effects, evidence references, and time.

Claim states distinguish proposed, unverified, partially verified, verified, refuted, and superseded records. Policy effects distinguish allow, warn, require approval, and block. These states make uncertainty and governance decisions explicit instead of collapsing them into a single confidence score. The current implementation does not require a non-empty evidence set for every verified claim and does not assess evidentiary sufficiency. This is an enforcement gap, not a guarantee supplied by the protocol. Figure 2 shows how the six record types reference one another across a full claim-to-outcome cycle.




Figure 2: The protocol’s recorded lifecycle. Solid arrows show synchronous references (a verification cites claims and evidence; an action proposal cites rationale claims; a policy decision governs a proposal; an outcome follows a decision). Dashed arrows show asynchronous feedback: a verification updates claim status, and an outcome’s observed effects may be registered as new evidence for future claims. The cycle is auditable, not automatic – each edge corresponds to an explicit recorded reference, not an implicit trust relationship.

4.2 Evidence ledger

The ledger supports JSON, SQLite, and PostgreSQL backends through a common interface. Collections are maintained for claims, evidence, verifications, actions, policy decisions, and outcomes.

Most records are appended. Claim verification performs a controlled in-place state transition and then rebuilds the affected collection chains. Each ordinary append computes

$$h_i = \text{SHA256}(h_{i-1} \parallel k \parallel \text{canonical}(r_i)),$$

where r_i is a record and k identifies its collection. Verification checks link continuity and recomputes each hash. Evidence payloads receive a SHA-256 content hash and may be signed with HMAC-SHA256 when a signing key is supplied.

Chain verification detects changes relative to the currently stored chain. This mechanism is *tamper evident*, not tamper proof: without an external anchor or protected signing key, it cannot distinguish malicious tampering from an authorized rebuild. An attacker with write access to both the records and all integrity keys can rebuild a consistent chain. Stronger deployment requires external key management, append-only storage, access control, backups, and independent audit anchors.

4.3 Transactional scoped ledger

Core 4.1.0-rc.6 adds a *scoped transactional* persistence layer for multi-tenant deployments. Business records are keyed by (*tenant_id*, *project_id*, *collection*, *canonical_key*) and stored with an explicit payload hash, chain metadata, and row version. The public API exposes `append_scoped`, `append_scoped_once`, `get_scoped`, `list_scoped`, `compare_and_set_scoped`, and `run_scoped_transaction`. A separate `verify_scoped_chain` operation recomputes hashes, chain continuity, and chain-head consistency without repairing data.

Write locks. Every mutating call inside a transaction must declare its locks up front. A *chain lock* covers an entire scoped collection and is required before appending to chain-linked GRDI collections; a *canonical lock* covers one business key within a collection. If a write touches a key without the matching declared lock, the engine raises `UndeclaredLockRef` rather than proceeding. Lock references are sorted in a global order (tenant, project, collection, kind, key) before acquisition to reduce deadlock risk across concurrent workers.

Atomic transactions. `run_scoped_transaction` accepts a callback that performs multiple scoped operations under one commit boundary. On the JSON backend, writes are staged in thread-local state and flushed on success; `transactional_mode=multiprocess` is rejected fail-closed because cross-process file locking is not provided. On SQLite, the engine uses `BEGIN IMMEDIATE` with scoped tables `phigraph_scoped_ledger` and `phigraph_chain_heads`. A callback failure or commit error rolls back staged state and clears transaction context.

Integrity verification. `verify_scoped_chain` groups rows by scope and collection, validates that internal columns match each record's `_chain` metadata, checks payload hashes, detects missing or orphan chain heads, and rejects sequence gaps. The check is *fail-closed*: any tamper or inconsistency raises `LedgerIntegrityError`. This behavior is *implemented* for JSON and SQLite in the current release; a PostgreSQL backend with the same semantics is *specified* but not yet shipped.

Figure 3 summarizes the transaction boundary and lock declaration requirement.




Figure 3: Scoped transactional writes require pre-declared locks and commit atomically on JSON (single-process) or SQLite (`BEGIN IMMEDIATE`). Integrity verification is a separate fail-closed read path.

4.4 Policy-gated runtime

The policy engine matches action type, permitted runtime modes, maximum risk, reversibility, and required approvals. No matching rule produces an implicit deny decision. For an action a , decision d , mode m , and executor x , the core execution condition is

$$\text{execute}(a) \iff (d.\text{effect} = \text{allow}) \wedge (m \in \{\text{copilot}, \text{guarded_auto}\}) \wedge (x \neq \emptyset).$$

Replay and shadow modes therefore generate simulated outcomes with `executed=false`. The packaged deployment configuration is stricter: it defaults to shadow-only operation, rejects attempts to disable that boundary, and disables real connectors. The core abstractions expose higher-authority modes for controlled future integration, but the supplied deployment does not grant arbitrary external execution.

4.5 API, SDK, and operations

The Python SDK can send tenant and project headers, API-key or bearer credentials, and idempotency keys. Authentication, PostgreSQL persistence, and signing are supported by lower-level Core components when explicitly configured. The packaged application does not currently wire every one of these options into the Core v3 router. Its Docker configuration must therefore be treated as a local evaluation setup unless protected by an external authentication and network boundary. Scope headers are application metadata and query filters; the present implementation does not provide end-to-end tenant isolation or same-scope referential-integrity guarantees.

4.6 Human-Assisted Verification (HAV v0.2)

HAV v0.2 is an additive module (`phigraph.hav`) that checks candidate textual output from an AI agent, such as a status summary or a release note, against an authoritative state supplied explicitly by the caller, and records the outcome through the existing Core protocol and ledger described in Section 4. HAV does not modify the ledger, policy engine, or runtime described above; it is a client of that protocol.

Pipeline. A verification request carries a candidate output string and an `AuthoritativeState`: a set of source-attributed evidence facts, or an explicit *unavailable* marker with a reason. A hybrid extractor combines fixed regular-expression patterns for structured, repository-style predicates (for example, test counts, CI or CodeQL status, and global success or production-readiness phrasing) with a second extractor that flags generic factual-looking substrings (percentages, dates, quantities, and source attributions) without asserting their truth. Every claim produced by the second extractor is marked `requires_external_grounding=true` and is not treated as verified evidence on its own. A verifier then compares each claim against the authoritative state by subject and predicate, and classifies it as supported, contradicted, unsupported, partially supported, or having insufficient evidence. Figure 4 traces this pipeline from candidate output to ledger records.




Figure 4: HAV v0.2 verification pipeline. The authoritative state is supplied by the caller, not retrieved by HAV, and merges into the verifier stage; an unavailable state is itself a valid input that the fail-closed policy (Figure 5) maps to a blocking verdict.

Fail-closed policy. A dedicated policy, `PHIGRAPH_HAV_FAIL_CLOSED_V1`, converts claim evaluations into one of five verdicts:

$$V(s, E) = \begin{cases} \text{SOURCE_UNAVAILABLE} & \text{if the authoritative state } s \text{ is unavailable} \\ \text{REJECT} & \text{if a critical claim in } E \text{ is contradicted} \\ \text{HUMAN_REVIEW} & \text{if a critical claim in } E \text{ lacks sufficient evidence} \\ \text{WARN} & \text{if any non-critical claim in } E \text{ is not supported} \\ \text{PASS} & \text{otherwise,} \end{cases}$$

evaluated in the listed order, where E is the set of claim evaluations for a request. Figure 5 renders this evaluation order as a decision flowchart. Verdicts map onto the existing PhiGraph policy-decision effects of Section 4 as shown in Table 1; in particular, an unavailable authoritative source blocks the action rather than defaulting to a permissive outcome.




Figure 5: Decision flowchart for the fail-closed policy $V(s, E)$. Branches are evaluated top to bottom; the first matching condition determines the verdict, so an unavailable source always blocks regardless of claim content, and a fully supported claim set only reaches PASS after every stricter condition fails.

Table 1: Mapping from HAV verdicts to PhiGraph Core policy effects.

HAV verdict	Core policy effect
PASS	allow
WARN	warn
REJECT	block
HUMAN_REVIEW	require_approval
SOURCE_UNAVAILABLE	block

Ledger integration and receipts. For each request, HAV registers the supplied evidence facts as **Evidence** records, registers each extracted claim as a **Claim** record, records a **Verification** per claim, registers an **ActionProposal** representing acceptance of the candidate output, and records a **PolicyDecision** carrying the mapped effect. A **HAVReceipt** bundling the verdict, claim evaluations, and a SHA-256 hash of the candidate output is itself stored as ledger evidence and, when a receipt-signing key is configured, signed with the same HMAC-SHA256 mechanism used elsewhere in Core. Tenant and project identifiers supplied with the request propagate to every record, consistent with the scope-tagging goal in Section 3.

Security posture. HAV never executes candidate output or any generated code; it only pattern-matches and compares strings. It does not persist a database file, and the packaged module contains no embedded credentials. The HTTP endpoints (`/v3/hav/verify`, `/v3/hav/factual/extract`, `/v3/hav/consistency`) accept an optional `PHIGRAPH_HAV_API_KEY`; when set, requests must present a matching `X-API-Key` header, compared using constant-time comparison, and are otherwise rejected. Multi-output consistency checking (shared-token overlap and conflicting status terms across several candidate outputs) is exposed only as an auxiliary diagnostic signal, not as a verification result on its own, and does not influence the policy decision above.

5 Evaluation

5.1 Research questions

The evaluation addresses three bounded questions:

- RQ1** Is the software version mechanically testable across its supported Python versions and packaging paths?
- RQ2** Does the current relational anomaly score produce useful rankings on one external network-flow benchmark relative to simple unsupervised baselines?
- RQ3** Does the HAV fail-closed policy produce the intended blocking or passing behavior on hand-constructed scenarios covering a contradicted critical claim, an unavailable authoritative source, and a supported numeric claim?

It does not test whether PhiGraph prevents all unsafe agent actions, proves causality, outperforms specialized state-of-the-art intrusion-detection systems, or measures HAV’s claim-extraction accuracy against a labeled claim corpus.

5.2 Software verification

The software tree contains 319 automated tests at commit `a5a7187` on `main`: 262 baseline tests covering protocol records, legacy and scoped ledger paths, policy evaluation, runtime modes, persistence backends, API and SDK boundaries, agent workflows, graph operations, benchmarks, and deployment configuration, plus 57 scoped transactional contract tests and 11 HAV-specific tests covering claim extraction, connectors, the fail-closed policy, the API surface, and API-key enforcement. Continuous integration is configured to install the declared extras and run the suite on Python 3.10, 3.11, 3.12, and 3.13; separate jobs build the Python distribution, install its wheel in an isolated environment, and build the Docker image. A local run on August 9, 2026 executed all 319 tests successfully. A separate synthetic benchmark attempt under local Python 3.14, which is outside the

CI matrix, stopped with an ARPACK eigensolver error and is excluded from the reported empirical results. Passing tests establish conformance to encoded cases, not correctness for all deployments.

5.3 HAV verification scenarios

To address RQ3, we exercised the running HAV API with three scenarios chosen to cover the fail-closed policy’s distinguishing branches, in addition to the 11 automated unit and integration tests. In the first, the candidate output asserts that all checks passed and the repository is production-ready while the supplied authoritative state marks CodeQL as failed and the release gate as blocked; the policy returned REJECT, mapped to a **block** policy decision recorded in the ledger. In the second, the authoritative source was marked unavailable; the policy returned SOURCE_UNAVAILABLE, also mapped to **block**, rather than defaulting to a permissive outcome. In the third, the candidate output correctly reports a numeric test count that matches the supplied evidence; the policy returned PASS. All three outcomes matched the behavior implied by the policy definition in Section 4.6. This is a confirmation of implemented decision logic on hand-constructed inputs, not a statistical evaluation of extraction accuracy, adversarial robustness, or coverage over naturally occurring agent output; the pattern-based extractor is not expected to generalize beyond the repository-status vocabulary encoded in its rules.

5.4 Transactional ledger integrity

To exercise the scoped API introduced in Section 4.3, the repository includes 57 contract tests (JSON and SQLite) covering idempotent append-once under concurrency, undeclared-lock rejection, thread-local JSON transaction isolation, compare-and-set on mutable Core collections without breaking append-only GRDI chains, commit-failure cleanup, and adversarial `verify_scoped_chain` cases (missing heads, orphan heads, sequence gaps, column tampering, and coordinated metadata manipulation). These tests establish *implemented* integrity behavior for the shipped backends; they do not prove security against a privileged attacker with signing keys and unrestricted storage access.

5.5 External anomaly-ranking experiment

The committed external-validation artifact uses the CIC-IDS2017 Friday afternoon DDoS flow file [9]. It contains 225,745 rows: 97,718 benign and 128,027 DDoS. With random seed 47, the protocol draws a 12,000-row benign-only reference set and evaluates a balanced test set of 6,000 benign and 6,000 DDoS flows. Preprocessing and variable-feature selection use only the benign reference. Methods are ranked continuously and evaluated at a top-10% alert budget.

The compared methods are robust multivariate deviation, Isolation Forest, Local Outlier Factor (LOF), and a PhiGraph relational score combining robust deviation, mean nearest-neighbor distance, and disagreement between their percentile ranks. Table 2 reproduces the committed result.

Table 2: CIC-IDS2017 Friday DDoS anomaly-ranking results. Runtime is reported by the committed validation artifact; its execution environment was not recorded.

Method	ROC-AUC	PR-AUC	P@10%	R@10%	F1@10%	Time (s)
Local Outlier Factor	0.9757	0.9506	0.9617	0.1923	0.3206	0.322
PhiGraph relational	0.7295	0.7316	0.8400	0.1680	0.2800	0.165
Isolation Forest	0.7403	0.6840	0.7817	0.1563	0.2606	0.321
Robust deviation	0.5771	0.5540	0.4117	0.0823	0.1372	0.006




Figure 6: PR-AUC and precision at a 10% alert budget for each method on the CIC-IDS2017 Friday DDoS experiment (same run as Table 2). Local Outlier Factor dominates on both metrics; the PhiGraph relational score ranks second, ahead of Isolation Forest and robust deviation.

LOF is clearly strongest on this experiment. PhiGraph is second by PR-AUC and precision at budget, but its ROC-AUC is slightly below Isolation Forest. This single fixed experiment establishes only that the implementation produces a non-degenerate ranking under the reported protocol. It does not support a claim of general or state-of-the-art superiority.

6 Threats to Validity and Limitations

Dataset validity. CIC-IDS2017 has documented problems in flow construction, feature generation, and labels [6, 4]. Results may change on corrected data. The machine-learning CSV also omits source and destination IP addresses, timestamps, users, devices, and processes. The evaluated graph is therefore a feature-similarity graph rather than the heterogeneous operational graph for which PhiGraph is designed.

Experimental validity. The external result uses one file, one attack family, one seed, and one alert budget. It reports no uncertainty intervals, ablation confidence intervals, or prospective evaluation. The benchmark is balanced and does not reproduce a production base rate. The PhiGraph score is batch-transductive because its component percentile ranks are calculated over the complete evaluation batch; scores can therefore change with batch composition and are not calibrated online anomaly scores. Additional chronological splits, repeated seeds, corrected datasets, heterogeneous event sources, and cost-sensitive metrics are required.

Software validity. Automated tests can encode incorrect assumptions and do not substitute for penetration testing, formal verification, independent replication, or operational incident analysis. The ledger’s hash chain detects accidental or unauthorized changes only under an appropriate trust and key-management model. JSON and SQLite are evaluation backends; production use requires hardened storage, secrets management, authorization, monitoring, and backup procedures.

HAV validity. HAV’s claim extractor is pattern- and rule-based, not a trained language or entailment model; its structured extractor recognizes a fixed vocabulary of repository-status phrasing in

English and Spanish, and its factual extractor flags generic percentage, date, quantity, and attribution patterns without grounding them. Recall and precision on naturally occurring agent output are therefore unmeasured, and the three scenarios in Section 5.3 are illustrative confirmations of policy logic, not a benchmark. Verification only compares extracted claims to caller-supplied evidence; HAV does not independently retrieve or authenticate that evidence, so an incorrect or manipulated authoritative state can still produce an incorrect verdict. Multi-output consistency checking is a lexical overlap heuristic and is explicitly documented as an auxiliary signal rather than proof of factual agreement.

Causal and governance claims. Graph connectivity, spectral structure, and anomaly scores do not identify causal effects without assumptions and suitable interventions. PhiGraph records candidate causal pathways but does not prove causality from observational data. Likewise, a recorded approval may still be mistaken or malicious. The platform improves inspectability and enforcement points; it does not guarantee ethical, legal, or correct outcomes.

7 Reproducibility and Availability

PhiGraph Core 4.0.0 is implemented in Python and distributed under the MIT License. Its development repository is maintained at <https://github.com/wcalmels/phigraph>. This paper is archived on Zenodo at <https://doi.org/10.5281/zenodo.21689514>. The paper source and references are licensed under CC BY 4.0.

The software test suite can be reproduced with:

```
python -m venv .venv
pip install -e ".[api,benchmark,dev,auth,app]"
pytest
```

The HAV-specific tests can be run in isolation with:

```
pytest tests/test_hav_core_integration.py tests/test_hav_api.py \
       tests/test_hav_v02_connectors.py \
       tests/test_hav_v02_factual_consistency.py \
       tests/test_hav_v02_benchmark.py \
       tests/test_hav_v02_api_security.py
```

The CIC-IDS2017 experiment requires the Friday DDoS CSV and is run with:

```
python scripts/validate_cicids2017.py /path/to/Friday...DDos.csv
```

The committed validation artifact records metrics and predictions but does not record the source-dataset checksum, exact dependency versions, hardware, operating system, or Git revision. It is therefore an auditable result artifact, not yet an exact reproducibility package. A final archive should preserve the dataset provenance and checksum, Python and dependency versions, seed, configuration, hardware description, and Git revision. The archival record is available at <https://doi.org/10.5281/zenodo.21689514>.

8 Ethical and Operational Considerations

PhiGraph can store sensitive evidence, identity metadata, and operational proposals. Deployers must minimize collection, define retention periods, encrypt transport and storage, separate tenant access, protect signing keys, and audit privileged operations. Security telemetry and network datasets may contain personal or confidential information and require an appropriate legal basis. The system should be evaluated in replay or shadow mode before any higher-authority integration. Human approval should remain mandatory where consequences are difficult to reverse or affect safety, rights, employment, finance, health, or critical infrastructure.

9 Conclusion

PhiGraph makes a specific architectural choice: intelligence does not imply authority. Claims must remain linked to evidence and verification; actions must remain subject to policy and explicit execution boundaries; outcomes must remain auditable. Version 4.0.0 demonstrates this design through a typed protocol, tamper-evident ledger, policy-gated runtime, scope-tagged records, and an automated test suite. HAV v0.2 extends the same design to a concrete instance of the problem: it treats agent-generated status claims as unverified until checked against an explicitly supplied authoritative state, and it fails closed, rather than open, when that state is unavailable. The external CIC-IDS2017 experiment shows a functional relational anomaly score but also illustrates why bounded interpretation is necessary: a conventional LOF baseline performs best, and the dataset cannot exercise PhiGraph’s full heterogeneous graph model. Future work should prioritize independent replication, corrected and chronological security datasets, multi-source operational graphs, a trained or retrieval-based claim extractor evaluated against a labeled corpus, policy-adversarial testing, cryptographically anchored audit storage, and prospective shadow deployments.

Author Contributions

Walter Calmels von Dem Knesebeck conceived the system, implemented the software, designed the evaluation, analyzed the results, and wrote the manuscript.

Conflict of Interest

The author is affiliated with TUCH Systems, the organization developing PhiGraph. This relationship may create a commercial interest in the system. The evaluation and limitations are reported to enable independent scrutiny.

License

This paper is licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0).

References

- [1] Peter W. Battaglia, Jessica B. Hamrick, Victor Bapst, Alvaro Sanchez-Gonzalez, Vinicius Zambaldi, Mateusz Malinowski, Andrea Tacchetti, David Raposo, Adam Santoro, Ryan Faulkner,

- et al. Relational inductive biases, deep learning, and graph networks. *arXiv preprint arXiv:1806.01261*, 2018.
- [2] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12):86–92, 2021.
 - [3] Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Delong Chen, Ho Shu Chan, Wenliang Dai, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38, 2023.
 - [4] Maxime Lanvin, Pierre-François Gimenez, Yufei Han, Frédéric Majorczyk, Ludovic Mé, and Éric Totel. Errors in the CICIDS2017 dataset and the significant differences in detection performances it makes. In *Risks and Security of Internet and Systems*, pages 18–33, 2022.
 - [5] Timothy Lebo, Satya Sahoo, and Deborah McGuinness. PROV-O: The PROV ontology. W3C recommendation, World Wide Web Consortium, 2013.
 - [6] Lisa Liu, Gints Engelen, Timothy Lynar, Daryl Essam, and Wouter Joosen. Error prevalence in NIDS datasets: A case study on CIC-IDS-2017 and CSE-CIC-IDS-2018. In *2022 IEEE Conference on Communications and Network Security*, pages 254–262, 2022.
 - [7] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency*, pages 220–229, 2019.
 - [8] OWASP Agentic Security Initiative. Agentic ai: Threats and mitigations. Technical report, OWASP Foundation, 2025.
 - [9] Iman Sharafaldin, Arash Habibi Lashkari, and Ali A. Ghorbani. Toward generating a new intrusion detection dataset and intrusion traffic characterization. In *Proceedings of the 4th International Conference on Information Systems Security and Privacy*, pages 108–116, 2018.
 - [10] Elham Tabassi. Artificial intelligence risk management framework (AI RMF 1.0). Technical Report NIST AI 100-1, National Institute of Standards and Technology, 2023.
 - [11] James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. FEVER: a large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 809–819. Association for Computational Linguistics, 2018.
 - [12] Santiago Torres-Arias, Hammad Afzali, Trishank Karthik Kuppusamy, Reza Curtmola, and Justin Cappos. in-toto: Providing farm-to-table guarantees for bits and bytes. In *28th USENIX Security Symposium*, pages 1393–1410. USENIX Association, 2019.